

# LLM Watermarking: Sequential Detection

# Motivating the Problem!

## Watermark:


A hidden signal embedded in text generated by a language model (LM) to trace its origin.

## Problem Setup:


Given an LM 'm' and a user prompt 'q', embed a watermark in the output 'y' such that:

1. **Distortion-Free:** Output quality is unchanged (i.e.,  $P(x) \approx$  original distribution).
2. **Model Agnostic:** Detection works without access to 'm' or 'q'.
3. **Robustness:** Watermark remains detectable even after adversarial modifications.

# Kirchenbauer - Red/Green List Watermarking



$$\text{Z-score} = \frac{(|s|_G - \gamma T)}{\sqrt{T\gamma(1-\gamma)}} = 4$$



Z-score >  $\tau$  (say 3)

**Detection**

**Overview of KGW**

## Solutions:

1. The technique is model/prompt agnostic and does not need the knowledge of model. (Although this a white-box approach)

## Problems:

1. Induces distortion into the output by changing probability distribution.
2. The technique is susceptible to substitution attacks.

# Kuditipudi - Statistical Watermarking


## Solutions:

1. Model/Prompt Agnostic
2. Since we do not manipulate the output logits of the LM, we mitigate the distortion.
3. The technique is robust to majority of attacks.

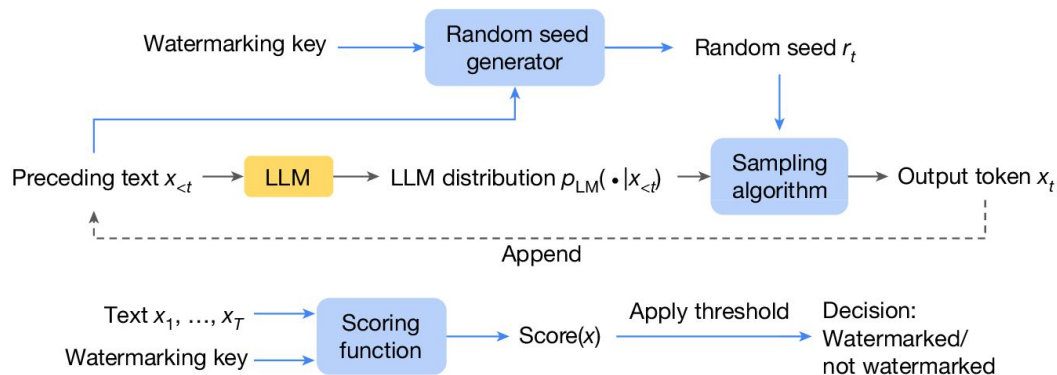
## Problems:

1. The algorithm has limitations on detection speeds, and is incapable of 'online' detection due to the design of the algorithm.

### LLM text generation



### Generative watermarking: text generation and watermark detection



# Statistical Setup of the Problem

**Objective:** Detect whether a given sequence of tokens  $Y_n = (Y_1, Y_2, \dots, Y_n)$  was generated by a watermarked language model.

**Setup:** Given a stream of observations:

Sequence of keys  $\xi = (\xi_1, \xi_2, \dots, \xi_n)$ ,

Output tokens  $Y_n = (Y_1, Y_2, \dots, Y_n)$

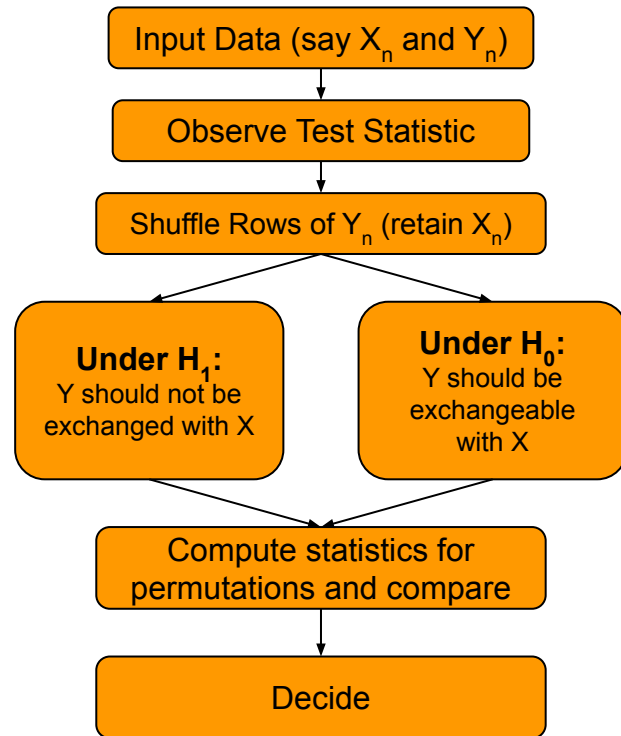
Decide between hypotheses:

$$H_0: \text{Text} \perp\!\!\!\perp \text{Key} \text{ vs. } H_a: \text{Text} \not\perp\!\!\!\perp \text{Key}$$

**Goal:** For  $\alpha \in (0, 1)$ , construct a level- $\alpha$  sequential test of power one,

→ Under  $H_0$ : continue forever w.p.  $\geq 1 - \alpha$

→ Under  $H_1$ : stop the test, and reject the null as soon as possible



For simplification, we modularise our watermarking scheme into 4 steps/algorithms:

→ **Step-1: Generating the watermark**

→ Step-2: Detecting the presence of watermark in a text

→ Step-3: Test statistic evaluating the misalignment between the keys and the text

---

**Algorithm 1:** Watermarked text generation (**generate**)

---

**Input** : watermark key sequence  $\xi \in \Xi^n$

**Params:** generation length  $m$ , language model  $p$ , decoder  $\Gamma$

**Output:** string  $y \in \mathcal{V}^m$

```
1 for  $i \in 1, \dots, m$  do  
2    $y_i \leftarrow \Gamma(\xi_{i \% n}, p(\cdot \mid y_{:i-1}))$   
3 return  $y$ 
```

---

For simplification, we modularise our watermarking scheme into 4 steps/algorithms:

→ Step-1: Generating the watermark

→ **Step-2: Detecting the presence of watermark in a text**

→ Step-3: Test statistic evaluating the misalignment between the keys and the text

---

**Algorithm 2:** Watermarked text detection (**detect**)

---

**Input** : string  $y \in \mathcal{V}^*$ , watermark key sequence  $\xi \in \Xi^n$

**Params:** test statistic  $\phi$ ; watermark key sequence distribution  $\nu \in \Delta(\Xi^n)$ ; resample size  $T$

**Output:** p-value  $\hat{p} \in [0, 1]$

```

1 for  $t \in 1, \dots, T$  do
2    $\xi^{(t)} \sim \nu$ 
3    $\phi_t \leftarrow \phi(y, \xi^{(t)})$ 
4  $\hat{p} \leftarrow \frac{1}{T+1} \left( 1 + \sum_{t=1}^T \mathbf{1}\{\phi_t \leq \phi(y, \xi)\} \right)$ 
5 return  $\hat{p}$ 

```

---

$\phi$  statistic values assuming text and key are independent (null)  
p-value = 0.003 (area to the left)

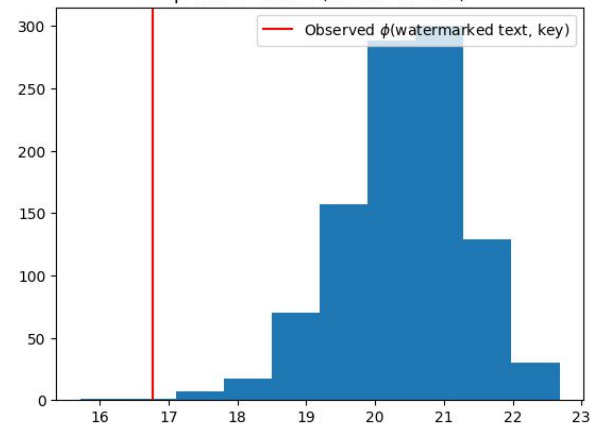


Fig: Distribution of test statistics (under the Null)

For simplification, we modularise our watermarking scheme into 4 steps/algorithms:

- Step-1: Generating the watermark
- Step-2: Detecting the presence of watermark in a text
- **Step-3: Test statistic evaluating the misalignment between the keys and the text**

---

**Algorithm 3:** Test statistic ( $\phi$ )

---

**Input** : string  $y \in \mathcal{V}^*$ , watermark key sequence  $\xi \in \Xi^n$

**Params:** alignment cost  $d$ , block size  $k$

**Output:** test statistic value  $\phi(y, \xi) \in \mathbb{R}$

```

1 for  $i \in 1, \dots, \text{len}(y) - k + 1$  do
2   for  $j \in 1, \dots, n$  do
3      $y^i \leftarrow \{y_{i+\ell}\}_{\ell=0}^{k-1}$ ,  $\xi^j \leftarrow \{\xi_{(j+\ell)\%n}\}_{\ell=0}^{k-1}$ 
4      $\hat{d}_{i,j} \leftarrow d(y^i, \xi^j)$ 
5 return  $\min_{i,j} \hat{d}_{i,j}$ 
```

---



# Proposed Solution

Main idea: Replace Kuditipudi et. al's permutation test with a *sequential* test

Sequential tests allow for gathering evidence against the null hypothesis in an online fashion and stop when it becomes significant. Compared to traditional (“batch”) tests:

- They often reach decisions much earlier (saving resources), and
- can make the same guarantees on their false positive rates

“Testing by betting” is an increasingly popular framework for designing sequential tests.

# Testing by Betting Framework

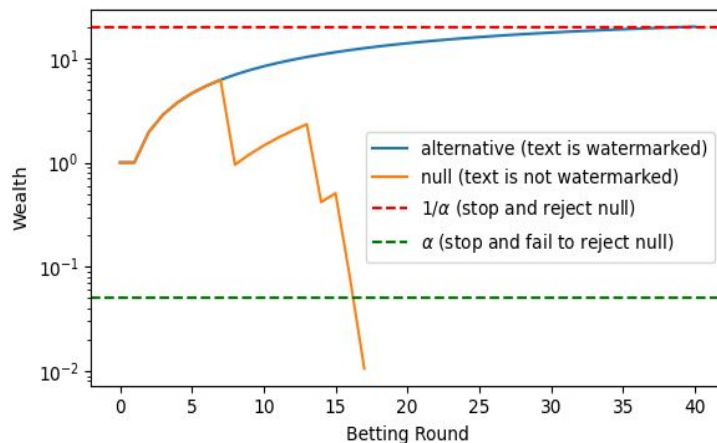
We track the wealth (a martingale by construction) of a gambler that bets against the null.

The betting function is designed such that the wealth (stochastic process):

- Is a martingale (remains constant in expectation) under the null
- Grows exponentially under the alternative

In our case,

- We are testing if the text is independent of the watermarking key
- Using a Monte Carlo permutation test, which computes  $T$  (expensive) test statistics
- Could before reaching  $T$  if we have enough evidence against the null?



→  $p\text{-value} = 1/(\text{wealth process})$

# Intuition and Guarantees of the Strategy

We focus on a log-optimal betting strategy designed specifically for hypothesis testing under a given alternative hypothesis.

Main Properties:

- **Log-Optimality:** Maximizes the expected log wealth under the considered alternative, ensuring statistically efficient use of evidence.
- **Finite-Time Guarantee:** Achieves zero resampling risk after a finite number of permutations — no need for infinite resampling to maintain validity.
- **Any-Time Valid:** The method maintains type-I error control at any stage, enabling real-time, sequential analysis without needing a fixed sample size.

For simplification, we modularise our watermarking scheme into 4 steps/algorithms:

- Step-1: Generating the watermark
- Step-2: Detecting the presence of watermark in a text
- Step-3: Test statistic evaluating the misalignment between the keys and the text
- **Step-4: Sequentializing the Hypothesis Test using a Sequential-MC Test**

---

**Algorithm 2':** Sequential Monte Carlo permutation test (`seq_mc_permutation_test`)

---

**Input:** tokens  $y \in \mathcal{V}^*$ , watermark key length  $n$ , block size  $k$ , test statistic function  $\phi$ , watermark key sequence  $\xi \in \Xi^n$ , threshold  $\alpha$ , slack parameter  $c$

**Output:** p-value estimate  $\hat{p} \in [0, 1]$ , runtime  $t \in \mathbb{N}$

```

1 begin
2    $W \leftarrow 1$ ;                                     // initial wealth
3    $L \leftarrow 0$ ;                                     // success count
4    $\phi_0 \leftarrow \phi(y, n, k)$ ;                     // flip sign of observed test statistic
5   for  $t = 1$  to  $T$  do
6      $\xi^{(t)} \sim \nu$ ;
7      $\phi_t \leftarrow \phi(y, \xi^{(t)})$ ;
8     if  $\phi_t \geq \phi_0$  then
9        $L \leftarrow L + 1$ ;                             // increment success count
10     $W \leftarrow \frac{1 - \text{BinomCDF}(L; t+1, c)}{c}$ ;           // update wealth using binomial tail
11    if  $W \geq \frac{1}{\alpha}$  or  $W < \alpha$  then
12      break;                                           // early stop
13   $\hat{p} \leftarrow 1 / \max(W, \epsilon)$ ;                   // final p-value estimate
14  return  $\hat{p}, t$ 

```

---

# Experiment setup

## Model and Dataset:

- OPT-1.3B — A 1.3 billion parameter open-source language model developed by Meta.
- C4 dataset— A large-scale English-language dataset curated for language modeling tasks.

## Watermark Generation Methods:

- ITS / ITS-edit: Inverse Transformed Sampling for watermarking and its sequential variant.
- EXP / EXP-edit: Exponential Sampling watermarking and its sequential variant.
- KGW-1.0: Kirchenbauer baselines for the sake of comparison.

## Evaluation Metrics:

- Permutation p-value: Used to test statistical dependence between generated tokens and watermarking mechanism.
- Number of Permutations: Reflects computational efficiency and convergence behavior of the test.

## Oracle Setup:

- We simulate an oracle setting where the watermark detection algorithm has access to the true distribution of the watermark signal under the null hypothesis (i.e., no watermarking).
- This setup enables us to isolate and evaluate the ideal performance of detection methods under best-case assumptions.

# Results: Watermarking the Midterm Report

## Midterm Report

As large language models (LLMs) continue to improve, traditional watermarking techniques—which previously depended on clear differences between machine-generated and human-written content—are becoming less reliable [9][3]. Alongside advancements in watermarking, more sophisticated detection techniques are also being developed [1][5][6].

Recently, watermarking strategies have explored statistical embedding and detection mechanisms. Notably, Kuditipudi [8] introduced a distortion-free watermarking approach coupled with a dependable detection method; however, its dependency on batch processing limits its practicality for real-time use.

To address this limitation, our project introduces a sequential watermark detection algorithm [2]. We propose an anytime-valid e-process/p-process framework that allows for real-time detection with early stopping, live access to test statistics, and reduced computational requirements.

**p-value:  $\sim 0.95 \Rightarrow$  Likely written by a human**

## Watermarked Report

As large language models (LLMs) continue to advance, the effectiveness of traditional watermarking methods—which often rely on observable distinctions between human-authored and machine-generated content—has diminished [9][3]. In parallel with improvements in watermarking strategies, researchers have also made progress in the development of more robust detection methodologies [1][5][6].


Recent approaches in watermarking have focused on statistical techniques for both embedding and detection. One such method, proposed by Kuditipudi [8], offers a distortion-free watermarking strategy alongside a reliable detection mechanism. However, this approach is limited by its reliance on batch processing, which constrains its applicability in real-time environments.

To overcome this limitation, the present work introduces a sequential watermark detection algorithm [2]. Our method leverages an anytime-valid e-process/p-process framework, which facilitates real-time detection through early stopping mechanisms, access to intermediate test statistics, and reduced computational demands.


**p-value: 0.0494  $\Rightarrow$  Likely written by LM**  
**\*(LM = OPT-1.3B here)**

# Results: Power & Null Rejection Rate

Power and Null Rejection Rate for c4 experiment without corruption  
text len (m)=80, key len (n)=256, # of texts (T)=200, alpha=0.05, c=0.04

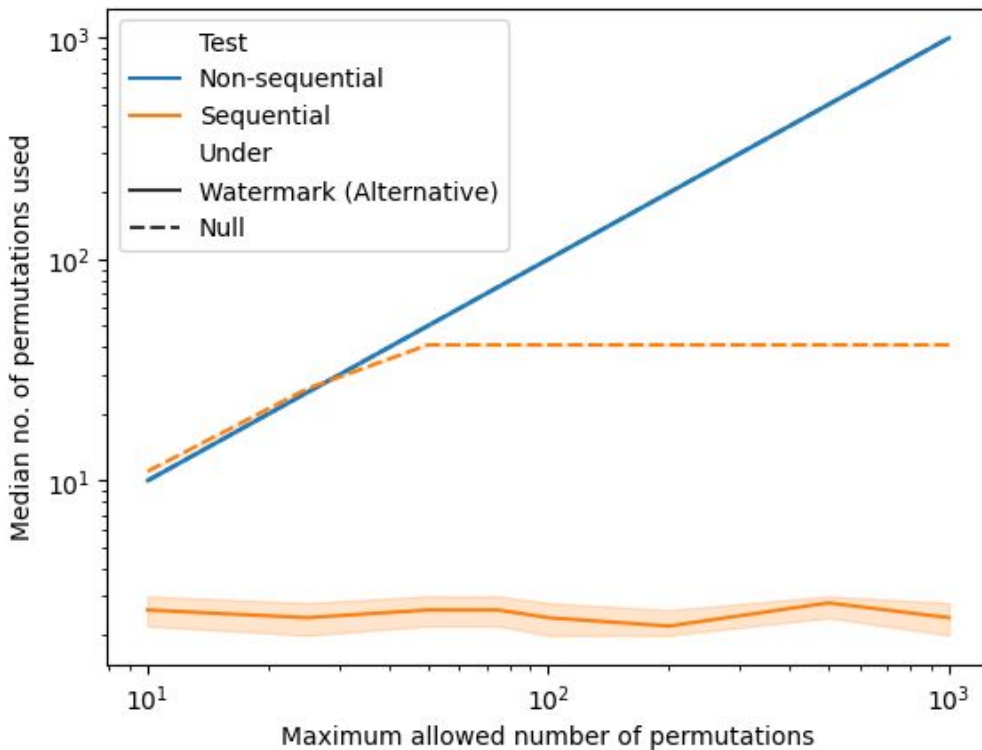


**Power =  $\mathbf{P}_{\text{Alternate}}$  (Reject Null) =  $1 - \beta$**   
a.k.a Type-2 error rate



**Null Rejection Rate =  $\mathbf{P}_{\text{Null}}$  (Reject Null) =  $\alpha$**   
a.k.a Type-1 error rate

# Results: Number of Permutations to Decision



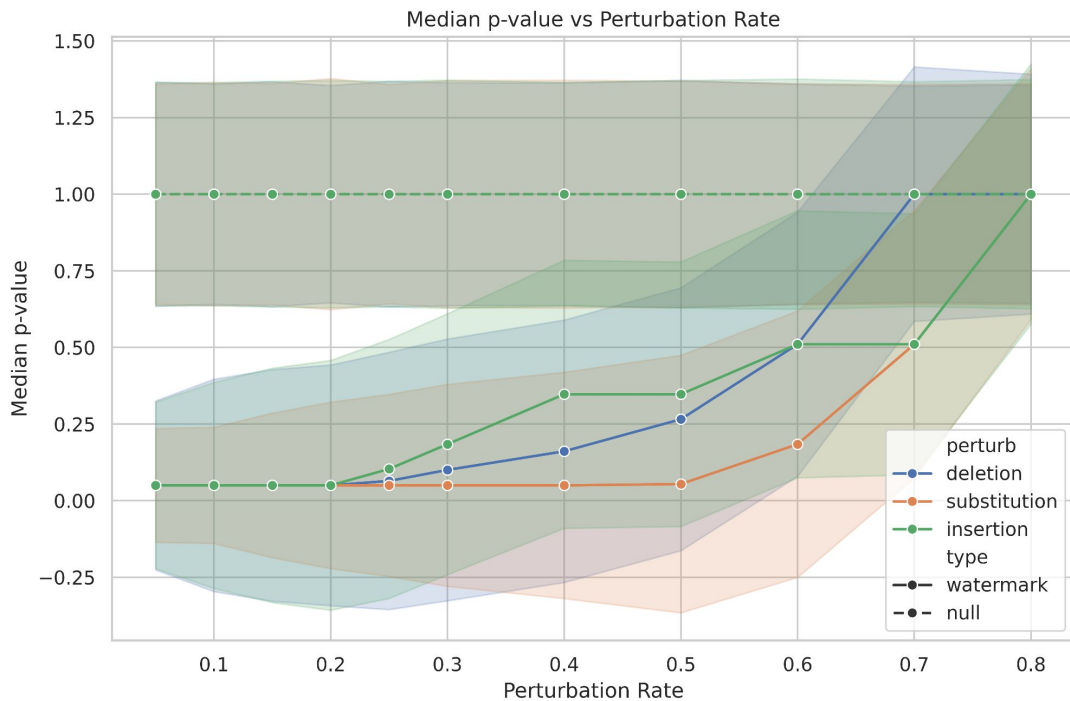
**Less Permutations  $\Rightarrow$  Less Time to Decision  $\Rightarrow$  Early Stopping  $\Rightarrow$  Less Computations**



# Experiment Setup: Testing Robustness

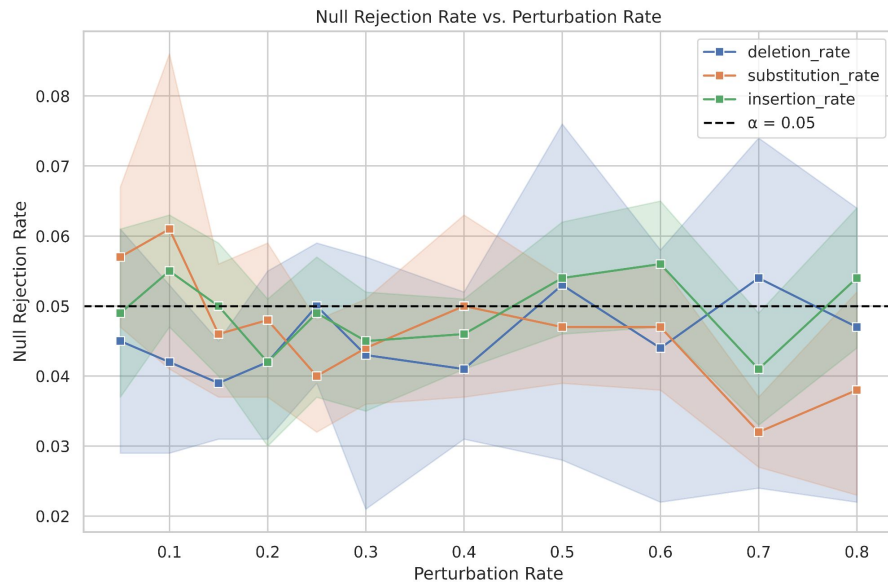
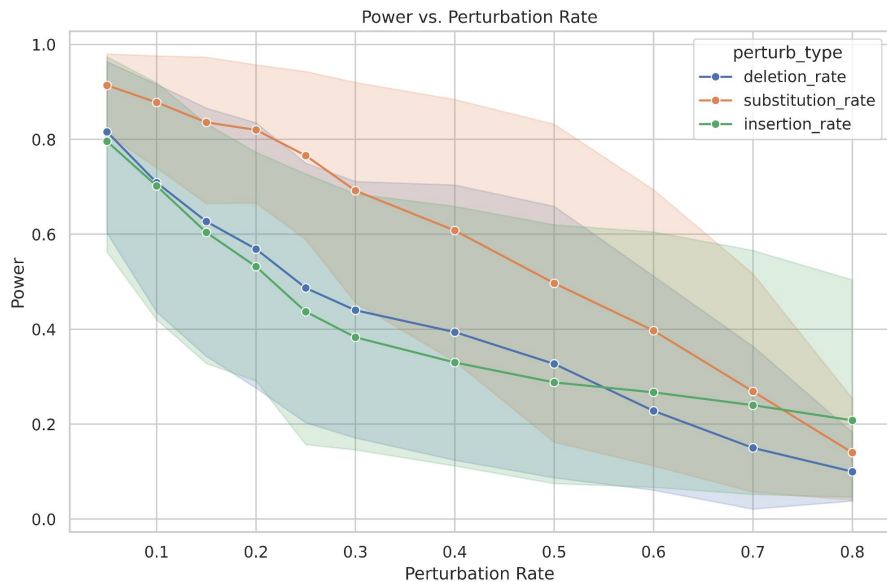
- Do  $T = 200$  permutations on a text of length  $m = 80$
- Remove/Insert/Substitute tokens from the generated output at rates ranging from 0.05 to 0.8 randomly.
- For each corruption rate, report:
  - ◆ the Average of Median p-values
  - ◆ Null Rejection Rate
  - ◆ Empirical Power
  - ◆ Median Permutation to Decision across Runs

# Results: Robustness of Watermark



**Median p-values remain low under perturbations  $\Rightarrow$  Watermark is Robust to attacks**

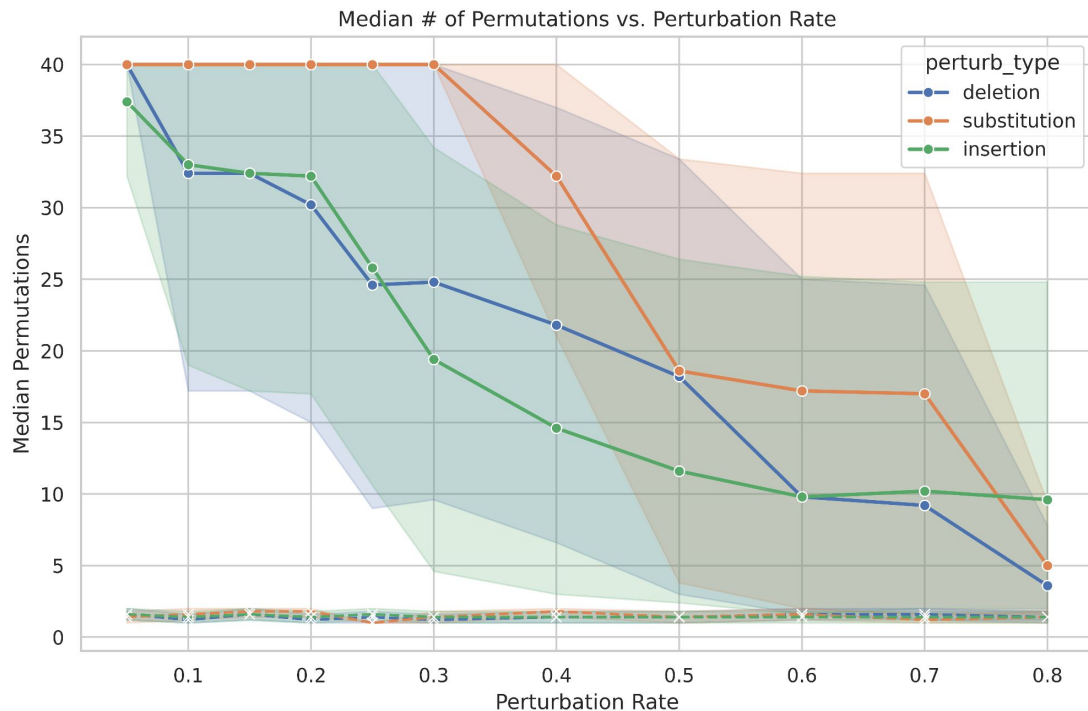
# Results: Power & Null Rejection Rate under Attack



**Power close to 1 + Null Rejection Rate below 0.05  $\Rightarrow$  Test remains valid under attack**

**\*Although the statistical confidence of the decision is impacted**

# Results: Number of Permutations to Decision under Attack



**Perturbation Increase  $\Rightarrow$  Watermark becomes weak  $\Rightarrow$  Algorithm initiates Early Stopping**  
**\*Still taking lesser time than permutations to make a decision**

# Analysis

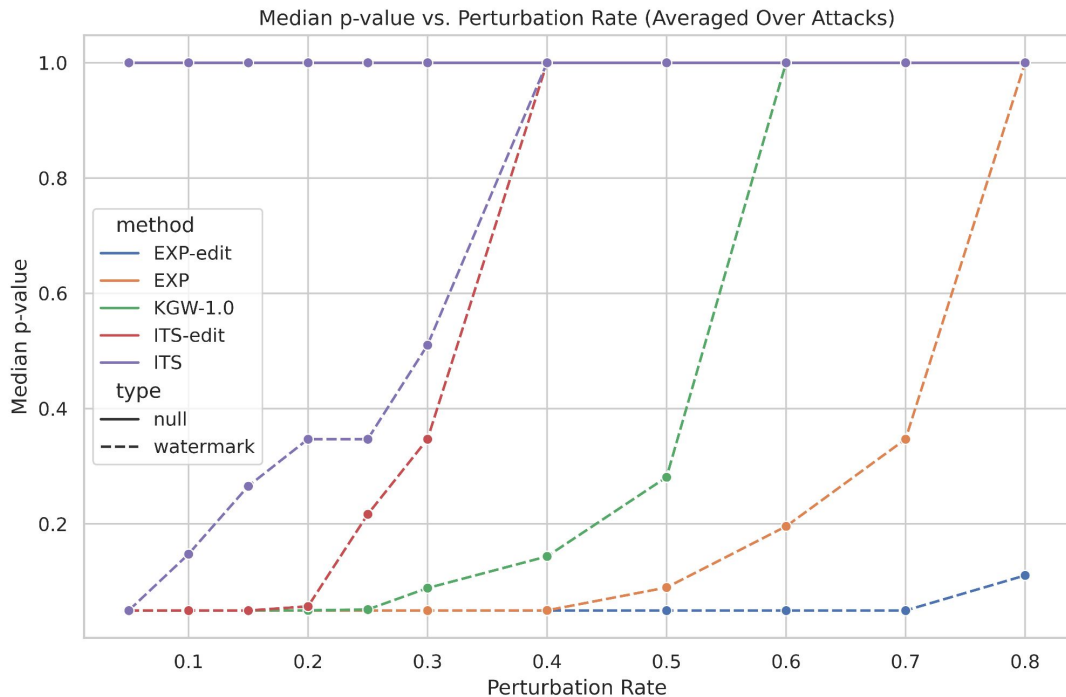
Substitution is most robust to corruption

Null rejection rate stays close to  $\alpha = 0.05$ , indicating that the false positive rate is well-controlled.

Require far less permutations

Require less to test null output


# Results: Comparing different generate algorithms




**EXP-edit method is most robust to attack outperforming EXP method!**

# Results: Performance under Attack

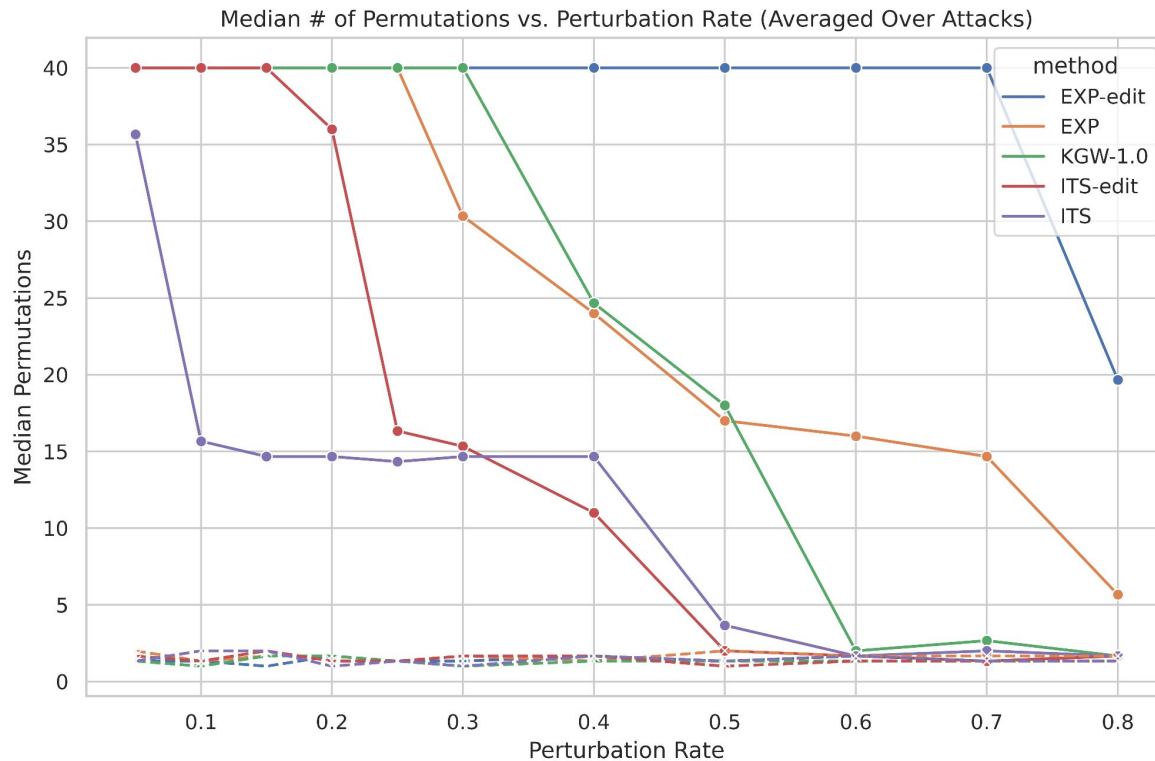
Power vs. Perturbation Rate (Averaged Over Attacks)



Null Rejection Rate vs. Perturbation Rate (Averaged Over Attacks)



# Results: Performance under Attack





# Conclusion

- **Robust Sequential Monte Carlo Test Implemented**: Developed a reliable and scalable SMC-based framework for watermark detection under real-world conditions.
- **Outperform the Permutation Tests**: Demonstrates significantly higher efficiency and greater robustness, especially in limited-sample or corrupted data settings.
- **High Statistical Power**: Power evaluations show strong and consistent detection capability, confirming the test's reliability across diverse scenarios.
- **Robust to Corruptions**: Maintains performance under token-level corruptions such as substitutions and deletions, making it practical for noisy or adversarial text.